


Regression and Classification

Dr. Tushar Sandhan

Content


- Curve fitting
- Modeling
 - Scalar linear model
- Loss function
- Multivariable regression
 - Vectorization
- Gradient descent optimization
- Polynomial regression
- Classification
 - Logistic regression

Curve fitting



Regression problem

- Given the inputs x , we need to estimate the targets t
- $\{(\mathbf{x}^{(i)}, t^{(i)})\}_{i=1}^N$
 - $\mathbf{x}^{(i)}$ inputs
 - $t^{(i)}$ targets
- Model:
 - $y = wx + b$



Modeling

- Given the inputs x , we need to estimate the targets t
- $\{(\mathbf{x}^{(i)}, t^{(i)})\}_{i=1}^N$
 - $\mathbf{x}^{(i)}$ inputs
 - t^i targets
- Model:
 - $y = wx + b$

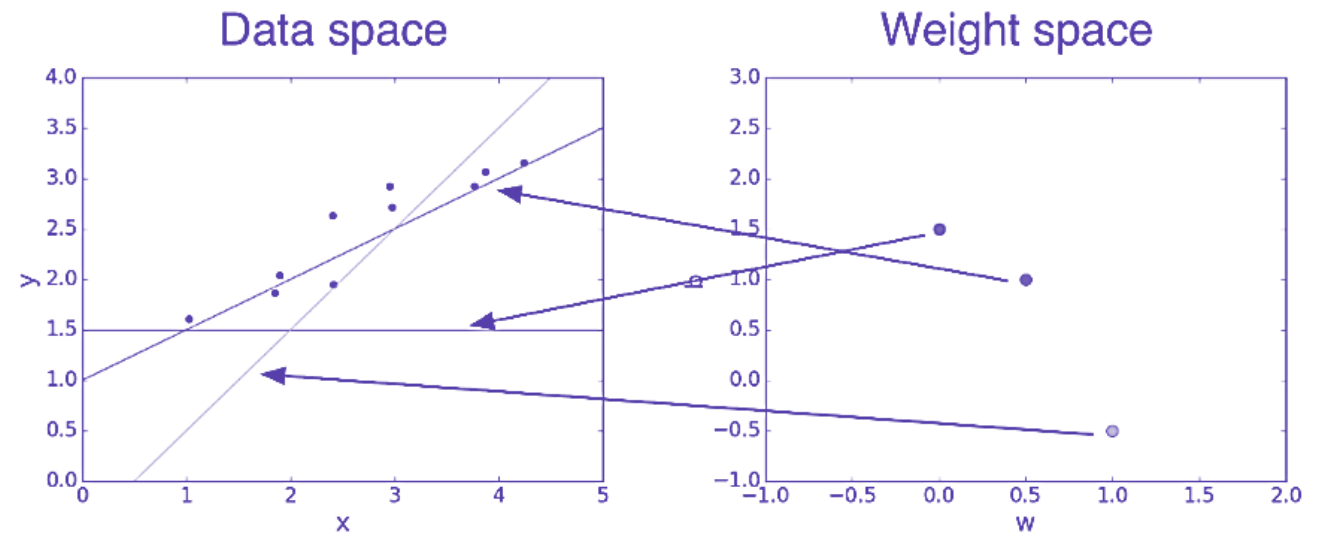
y is the **prediction**

w is the **weight**

b is the **bias**


w and b together are the **parameters**

Settings of the parameters are called **hypotheses**



Loss function

- Loss function: sq. error
- Residual
 - $y - t$
- Cost function:



$$\mathcal{L}(y, t) = \frac{1}{2}(y - t)^2$$

$$\mathcal{J}(w, b) = \frac{1}{N} \sum_{i=1}^N \left(y^{(i)} - t^{(i)} \right)^2$$

- Model:

- $y = wx + b$

$$= \frac{1}{N} \sum_{i=1}^N \left(wx^{(i)} + b - t^{(i)} \right)^2$$

Multivariate regression

- MISO
 - Not much different than single input (scalar) case
- MIMO
 - targets also multiple
 - Multilayer perceptrons

- MISO:

$$y = \sum_j w_j x_j + b$$

Solution (1)

Partial derivatives: derivatives of a multivariate function with respect to one of its arguments.

$$y = \sum_j w_j x_j + b$$

$$\frac{\partial}{\partial x_1} f(x_1, x_2) = \lim_{h \rightarrow 0} \frac{f(x_1 + h, x_2) - f(x_1, x_2)}{h}$$

To compute, take the single variable derivatives, pretending the other arguments are constant.

Example: partial derivatives of the prediction y

$$\begin{aligned} \frac{\partial y}{\partial w_j} &= \frac{\partial}{\partial w_j} \left[\sum_{j'} w_{j'} x_{j'} + b \right] \\ &= x_j \end{aligned}$$

$$\begin{aligned} \frac{\partial y}{\partial b} &= \frac{\partial}{\partial b} \left[\sum_{j'} w_{j'} x_{j'} + b \right] \\ &= 1 \end{aligned}$$

Solution (1)

Chain rule for derivatives:

$$\begin{aligned}\frac{\partial \mathcal{L}}{\partial w_j} &= \frac{d\mathcal{L}}{dy} \frac{\partial y}{\partial w_j} \\ &= \frac{d}{dy} \left[\frac{1}{2} (y - t)^2 \right] \cdot x_j \\ &= (y - t) x_j \\ \frac{\partial \mathcal{L}}{\partial b} &= y - t\end{aligned}$$

Cost derivatives (average over data points):

$$\begin{aligned}\frac{\partial \mathcal{J}}{\partial w_j} &= \frac{1}{N} \sum_{i=1}^N (y^{(i)} - t^{(i)}) x_j^{(i)} \\ \frac{\partial \mathcal{J}}{\partial b} &= \frac{1}{N} \sum_{i=1}^N y^{(i)} - t^{(i)}\end{aligned}$$

The minimum must occur at a point where the partial derivatives are zero.

$$\frac{\partial \mathcal{J}}{\partial w_j} = 0 \quad \frac{\partial \mathcal{J}}{\partial b} = 0.$$

Optimal weights:

$$\mathbf{w} = (\mathbf{X}^\top \mathbf{X})^{-1} \mathbf{X}^\top \mathbf{t}$$

Solution (2)

- Gradient descent
- Iterative algo
 - apply update repeatedly until convergence

Observe:

- if $\partial \mathcal{J} / \partial w_j > 0$, then increasing w_j increases \mathcal{J} .
- if $\partial \mathcal{J} / \partial w_j < 0$, then increasing w_j decreases \mathcal{J} .

Solution (2): Gradient descent

- Gradient descent
- Iterative algo
 - apply update repeatedly until convergence

The following update decreases the cost function:

$$\begin{aligned}w_j &\leftarrow w_j - \alpha \frac{\partial \mathcal{J}}{\partial w_j} \\&= w_j - \frac{\alpha}{N} \sum_{i=1}^N (y^{(i)} - t^{(i)}) x_j^{(i)}\end{aligned}$$

α is a **learning rate**. The larger it is, the faster \mathbf{w} changes.

- We'll see later how to tune the learning rate, but values are typically small, e.g. 0.01 or 0.0001

Update rule vector form

This gets its name from the **gradient**:



$$\frac{\partial \mathcal{J}}{\partial \mathbf{w}} = \begin{pmatrix} \frac{\partial \mathcal{J}}{\partial w_1} \\ \vdots \\ \frac{\partial \mathcal{J}}{\partial w_D} \end{pmatrix}$$

- This is the direction of fastest increase in \mathcal{J} .

Update rule in vector form:

$$\begin{aligned} \mathbf{w} &\leftarrow \mathbf{w} - \alpha \frac{\partial \mathcal{J}}{\partial \mathbf{w}} \\ &= \mathbf{w} - \frac{\alpha}{N} \sum_{i=1}^N (y^{(i)} - t^{(i)}) \mathbf{x}^{(i)} \end{aligned}$$

Hence, gradient descent updates the weights in the direction of fastest *decrease*.



Polynomial regression

- Regression using higher degree polynomials
- Data distributed non-linear way


$$y = w_3x^3 + w_2x^2 + w_1x + w_0$$

Define the **feature map**

$$\psi(x) = \begin{pmatrix} 1 \\ x \\ x^2 \\ x^3 \end{pmatrix}$$


Polynomial regression model:

$$y = \mathbf{w}^T \psi(x)$$




Hyper-parameter: degree of polynomial


$$y = w_0$$




$$y = w_0 + w_1x$$



$$y = w_0 + w_1x + w_2x^2 + w_3x^3$$




$$y = w_0 + w_1x + w_2x^2 + w_3x^3 + \dots + w_9x^9$$




Classification: logistic regression

- Target distribution is quantized (class labels/probabilities)



$y(\mathbf{x}) = \sigma(\mathbf{w}^T \mathbf{x} + w_0)$
where the sigmoid is defined as

$$\sigma(z) = \frac{1}{1 + \exp(-z)}$$



Classification

- Logistic regression can be considered as last layer of neural network
- Inputs are x^i , weights are w
- Sigmoid function is non-linear activation
- Similar manner, we compare prediction error and minimize loss by updating weights

$$\begin{aligned} J(\theta) &= \sum_{n=1}^N \mathcal{L}(h_{\theta}(\mathbf{x}_n), y_n) \\ &= \sum_{n=1}^N -\left\{ y_n \log h_{\theta}(\mathbf{x}_n) + (1 - y_n) \log(1 - h_{\theta}(\mathbf{x}_n)) \right\} \end{aligned}$$

Thank you