

Demo python programming

Dr. Tushar Sandhan

▼ Setup

Import TensorFlow and other necessary libraries:

```
[ ] import matplotlib.pyplot as plt
    import numpy as np
    import PIL
    import tensorflow as tf

    from tensorflow import keras
    from tensorflow.keras import layers
    from tensorflow.keras.models import Sequential
```

```
[ ] import pathlib
    dataset_url = "https://storage.googleapis.com/download.tensorflow.org/example_images/flower_photos.tgz"
    data_dir = tf.keras.utils.get_file('flower_photos', origin=dataset_url, untar=True)
    data_dir = pathlib.Path(data_dir)
```

Downloading data from https://storage.googleapis.com/download.tensorflow.org/example_images/flower_photos.tgz
228813984/228813984 [=====] - 3s 0us/step

After downloading, you should now have a copy of the dataset available. There are 3,670 total images:

```
[ ] image_count = len(list(data_dir.glob('*/*.jpg')))
    print(image_count)
```

3670

Here are some roses:

```
[ ] roses = list(data_dir.glob('roses/*'))  
PIL.Image.open(str(roses[0]))
```



```
[ ] PIL.Image.open(str(roses[1]))
```



And some tulips:

```
[ ] tulips = list(data_dir.glob('tulips/*'))  
PIL.Image.open(str(tulips[0]))
```



```
[ ] PIL.Image.open(str(tulips[1]))
```



Define some parameters for the loader:

```
[ ] batch_size = 32  
    img_height = 180  
    img_width = 180
```

```
[ ] train_ds = tf.keras.utils.image_dataset_from_directory(  
    data_dir,  
    validation_split=0.2,  
    subset="training",  
    seed=123,  
    image_size=(img_height, img_width),  
    batch_size=batch_size)
```

```
Found 3670 files belonging to 5 classes.  
Using 2936 files for training.
```

```
[ ] val_ds = tf.keras.utils.image_dataset_from_directory(
    data_dir,
    validation_split=0.2,
    subset="validation",
    seed=123,
    image_size=(img_height, img_width),
    batch_size=batch_size)
```

```
Found 3670 files belonging to 5 classes.
Using 734 files for validation.
```

You can find the class names in the `class_names` attribute on these datasets. These correspond to the directory names in alphabetical order.

```
[ ] class_names = train_ds.class_names
    print(class_names)

['daisy', 'dandelion', 'roses', 'sunflowers', 'tulips']
```

Visualize the data

Here are the first nine images from the training dataset:

```
[ ] import matplotlib.pyplot as plt

plt.figure(figsize=(10, 10))
for images, labels in train_ds.take(1):
    for i in range(9):
        ax = plt.subplot(3, 3, i + 1)
        plt.imshow(images[i].numpy().astype("uint8"))
        plt.title(class_names[labels[i]])
        plt.axis("off")
```

roses



dandelion



tulips



sunflowers



dandelion



roses



dandelion



roses



tulips



You will pass these datasets to the Keras `Model.fit` method for training later in this tutorial. If you like, you can also manually iterate over the dataset and retrieve batches of images:

```
[ ] for image_batch, labels_batch in train_ds:  
    print(image_batch.shape)  
    print(labels_batch.shape)  
    break
```

```
(32, 180, 180, 3)  
(32,)
```

▼ Configure the dataset for performance

Make sure to use buffered prefetching, so you can yield data from disk without having I/O become blocking.

```
▶ AUTOTUNE = tf.data.AUTOTUNE  
  
train_ds = train_ds.cache().shuffle(1000).prefetch(buffer_size=AUTOTUNE)  
val_ds = val_ds.cache().prefetch(buffer_size=AUTOTUNE)
```

▼ Standardize the data

The RGB channel values are in the `[0, 255]` range. This is not ideal for a neural network; in general you should seek to make your input values small.

Here, you will standardize values to be in the `[0, 1]` range by using `tf.keras.layers.Rescaling`:

```
[ ] normalization_layer = layers.Rescaling(1./255)
```

There are two ways to use this layer. You can apply it to the dataset by calling `Dataset.map`:

```
[ ] normalized_ds = train_ds.map(lambda x, y: (normalization_layer(x), y))
    image_batch, labels_batch = next(iter(normalized_ds))
    first_image = image_batch[0]
    # Notice the pixel values are now in `[0,1]`.
    print(np.min(first_image), np.max(first_image))
```

```
0.0 1.0
```

▼ A basic CNN model

Create the model

```
[ ] num_classes = len(class_names)

model = Sequential([
    layers.Rescaling(1./255, input_shape=(img_height, img_width, 3)),
    layers.Conv2D(16, 3, padding='same', activation='relu'),
    layers.MaxPooling2D(),
    layers.Conv2D(32, 3, padding='same', activation='relu'),
    layers.MaxPooling2D(),
    layers.Conv2D(64, 3, padding='same', activation='relu'),
    layers.MaxPooling2D(),
    layers.Flatten(),
    layers.Dense(128, activation='relu'),
    layers.Dense(num_classes)
])
```

▼ Compile the model

```
[ ] model.compile(optimizer='adam',  
                  loss=tf.keras.losses.SparseCategoricalCrossentropy(from_logits=True),  
                  metrics=['accuracy'])
```

▼ Model summary

View all the layers of the network using the Keras `Model.summary` method:

```
[ ] model.summary()
```

Model: "sequential_1"

Layer (type)	Output Shape	Param #
rescaling_2 (Rescaling)	(None, 180, 180, 3)	0
conv2d_3 (Conv2D)	(None, 180, 180, 16)	448
max_pooling2d_3 (MaxPooling 2D)	(None, 90, 90, 16)	0
conv2d_4 (Conv2D)	(None, 90, 90, 32)	4640
max_pooling2d_4 (MaxPooling 2D)	(None, 45, 45, 32)	0
conv2d_5 (Conv2D)	(None, 45, 45, 64)	18496
max_pooling2d_5 (MaxPooling 2D)	(None, 22, 22, 64)	0
flatten_1 (Flatten)	(None, 30976)	0
dense_2 (Dense)	(None, 128)	3965056
dense_3 (Dense)	(None, 5)	645
Total params: 3,989,285		
Trainable params: 3,989,285		
Non-trainable params: 0		

▼ Train the model

Train the model for 10 epochs with the Keras `Model.fit` method:

```
[ ] epochs=10
    history = model.fit(
        train_ds,
        validation_data=val_ds,
        epochs=epochs
    )
```

```
Epoch 1/10
92/92 [=====] - 11s 37ms/step - loss: 1.3572 - accuracy: 0.4012 - val_loss: 1.1089 - val_accuracy: 0.5681
Epoch 2/10
92/92 [=====] - 2s 23ms/step - loss: 1.0010 - accuracy: 0.6046 - val_loss: 0.9536 - val_accuracy: 0.5940
Epoch 3/10
92/92 [=====] - 2s 23ms/step - loss: 0.8055 - accuracy: 0.7027 - val_loss: 0.9053 - val_accuracy: 0.6594
Epoch 4/10
92/92 [=====] - 2s 24ms/step - loss: 0.5934 - accuracy: 0.7854 - val_loss: 0.9686 - val_accuracy: 0.6526
Epoch 5/10
92/92 [=====] - 2s 24ms/step - loss: 0.3918 - accuracy: 0.8658 - val_loss: 1.0422 - val_accuracy: 0.6594
Epoch 6/10
92/92 [=====] - 2s 25ms/step - loss: 0.2248 - accuracy: 0.9244 - val_loss: 1.1349 - val_accuracy: 0.6594
Epoch 7/10
92/92 [=====] - 2s 23ms/step - loss: 0.1493 - accuracy: 0.9554 - val_loss: 1.3704 - val_accuracy: 0.6608
Epoch 8/10
92/92 [=====] - 2s 23ms/step - loss: 0.1043 - accuracy: 0.9687 - val_loss: 1.5848 - val_accuracy: 0.6730
Epoch 9/10
92/92 [=====] - 2s 23ms/step - loss: 0.0328 - accuracy: 0.9939 - val_loss: 1.9913 - val_accuracy: 0.6485
Epoch 10/10
92/92 [=====] - 2s 25ms/step - loss: 0.0170 - accuracy: 0.9973 - val_loss: 2.1011 - val_accuracy: 0.6553
```

Visualize training results

Create plots of the loss and accuracy on the training and validation sets:


```
[ ] acc = history.history['accuracy']
    val_acc = history.history['val_accuracy']

    loss = history.history['loss']
    val_loss = history.history['val_loss']

    epochs_range = range(epochs)

    plt.figure(figsize=(8, 8))
    plt.subplot(1, 2, 1)
    plt.plot(epochs_range, acc, label='Training Accuracy')
    plt.plot(epochs_range, val_acc, label='Validation Accuracy')
    plt.legend(loc='lower right')
    plt.title('Training and Validation Accuracy')

    plt.subplot(1, 2, 2)
    plt.plot(epochs_range, loss, label='Training Loss')
    plt.plot(epochs_range, val_loss, label='Validation Loss')
    plt.legend(loc='upper right')
    plt.title('Training and Validation Loss')
    plt.show()
```



State of the Art-CNN model Resnet50

```
[ ] from tensorflow.keras.models import Model
    from tensorflow.keras.layers import Conv2D, Flatten, Dense, MaxPool2D, BatchNormalization, GlobalAveragePooling2D
    from tensorflow.keras.applications.resnet50 import preprocess_input, decode_predictions
    from tensorflow.keras.preprocessing.image import ImageDataGenerator, load_img
    from tensorflow.keras.applications.resnet50 import ResNet50
    from tensorflow.keras.preprocessing import image
```

```
[ ] base_model = ResNet50(include_top=False,weights= 'imagenet')
    x = base_model.output
    x = GlobalAveragePooling2D()(x)
    x = Dense(1024,activation='relu')(x)
    predictions = Dense(5,activation='softmax')(x)
    model = Model(inputs=base_model.input,outputs=predictions)

    for layer in base_model.layers:
        layer.trainable = False

    model.compile(optimizer='adam',loss='categorical_crossentropy',metrics = ['accuracy'])
```

Downloading data from https://storage.googleapis.com/tensorflow/keras-applications/resnet/resnet50_weights_tf_dim_ordering_tf_kernels_notop.h5
94765736/94765736 [=====] - 0s 0us/step

```
[ ] model.compile(optimizer='adam',
                  loss=tf.keras.losses.SparseCategoricalCrossentropy(from_logits=True),
                  metrics=['accuracy'])
```

Model summary

View all the layers of the network using the Keras Model.summary method:

```
[ ] model.summary()

conv2_block3_1_bn (BatchNormal (None, None, None, 256      ['conv2_block3_1_conv[0][0]']
ization)          64)

conv2_block3_1_relu (Activatio (None, None, None, 0      ['conv2_block3_1_bn[0][0]']
n)                64)

conv2_block3_2_conv (Conv2D)    (None, None, None, 36928   ['conv2_block3_1_relu[0][0]']
64)

conv2_block3_2_bn (BatchNormal (None, None, None, 256      ['conv2_block3_2_conv[0][0]']
ization)          64)

conv2_block3_2_relu (Activatio (None, None, None, 0      ['conv2_block3_2_bn[0][0]']
n)                64)

conv2_block3_3_conv (Conv2D)    (None, None, None, 16640   ['conv2_block3_2_relu[0][0]']
256)

conv2_block3_3_bn (BatchNormal (None, None, None, 1024     ['conv2_block3_3_conv[0][0]']
ization)          256)

      ■
      ■
      ■

global_average_pooling2d (Glob (None, 2048)      0      ['conv5_block3_out[0][0]']
alAveragePooling2D)

dense_4 (Dense)           (None, 1024)      2098176  ['global_average_pooling2d[0][0]'
]

dense_5 (Dense)           (None, 5)          5125     ['dense_4[0][0]']

=====
Total params: 25,691,013
Trainable params: 2,103,301
Non-trainable params: 23,587,712
```



```
epochs=10
history = model.fit(
    train_ds,
    validation_data=val_ds,
    epochs=epochs
)
```

Epoch 1/10

/usr/local/lib/python3.8/dist-packages/tensorflow/python/util/dispatch.py:1082: UserWarning: ``sparse_categorical_crossentropy`` received
return dispatch_target(*args, **kwargs)

92/92 [=====] - 14s 110ms/step - loss: 0.7490 - accuracy: 0.7793 - val_loss: 0.3433 - val_accuracy: 0.8828

Epoch 2/10

92/92 [=====] - 8s 89ms/step - loss: 0.2634 - accuracy: 0.9077 - val_loss: 0.3493 - val_accuracy: 0.8815

Epoch 3/10

92/92 [=====] - 8s 88ms/step - loss: 0.1460 - accuracy: 0.9482 - val_loss: 0.4676 - val_accuracy: 0.8447

Epoch 4/10

92/92 [=====] - 8s 89ms/step - loss: 0.0860 - accuracy: 0.9741 - val_loss: 0.4188 - val_accuracy: 0.8719

Epoch 5/10

92/92 [=====] - 8s 90ms/step - loss: 0.0444 - accuracy: 0.9905 - val_loss: 0.3737 - val_accuracy: 0.8815

Epoch 6/10

92/92 [=====] - 8s 91ms/step - loss: 0.0207 - accuracy: 0.9966 - val_loss: 0.4846 - val_accuracy: 0.8583

Epoch 7/10

92/92 [=====] - 9s 102ms/step - loss: 0.0175 - accuracy: 0.9976 - val_loss: 0.3954 - val_accuracy: 0.9005

Epoch 8/10

92/92 [=====] - 8s 91ms/step - loss: 0.0056 - accuracy: 1.0000 - val_loss: 0.4361 - val_accuracy: 0.8828

Epoch 9/10

92/92 [=====] - 8s 91ms/step - loss: 0.0034 - accuracy: 1.0000 - val_loss: 0.4233 - val_accuracy: 0.8951

Epoch 10/10

92/92 [=====] - 8s 91ms/step - loss: 0.0019 - accuracy: 1.0000 - val_loss: 0.4222 - val_accuracy: 0.8910


```
[ ] acc = history.history['accuracy']
    val_acc = history.history['val_accuracy']

    loss = history.history['loss']
    val_loss = history.history['val_loss']

    epochs_range = range(epochs)

    plt.figure(figsize=(8, 8))
    plt.subplot(1, 2, 1)
    plt.plot(epochs_range, acc, label='Training Accuracy')
    plt.plot(epochs_range, val_acc, label='Validation Accuracy')
    plt.legend(loc='lower right')
    plt.title('Training and Validation Accuracy')

    plt.subplot(1, 2, 2)
    plt.plot(epochs_range, loss, label='Training Loss')
    plt.plot(epochs_range, val_loss, label='Validation Loss')
    plt.legend(loc='upper right')
    plt.title('Training and Validation Loss')
    plt.show()
```



Thank you